

ml_project

March 2, 2020

```
In [38]: import pandas as pd
import numpy as np
import scipy.linalg as la
import scipy.stats as stats
import statsmodels.api as sm
from sklearn import linear_model, model_selection, metrics
import sklearn
import plotly.graph_objs as go
import matplotlib.pyplot as plt
from sklearn.datasets import load_breast_cancer, load_digits, load_diabetes
from sklearn.model_selection import train_test_split
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA
from sklearn.neighbors import KNeighborsClassifier as KNC
from sklearn.discriminant_analysis import QuadraticDiscriminantAnalysis as GDA
from sklearn.metrics import roc_auc_score as RAS
from sklearn.metrics import roc_curve as ROC
import seaborn as sns
from sklearn.manifold import TSNE
import umap
```

```
jupyter nbconvert --to latex "Race and Incarceration in America.ipynb"
--TagRemovePreprocessor.remove_input_tags='{"hide_input"}'; pdflatex "Race and
Incarceration in America"
```

1 Introduction

Last semester I took a look a data set containing the information of more than 7.5 million individuals that have been processed by the criminal justice system. I found that racial minorities were more likely to receive extreme sentences, agreeing with existing research around bias in the criminal justice system. In this project I will be exploring the data from a machine learning perspective. My goal is to determine if this data can be classified in such a way that is predictive of race.

2 Possible questions

can we use ml techniques to correctly classify this data? which ones fail and why? can we create a predictive model for sentence lengths? should sentencing be offloaded to a ml algorithm.

```

In [14]: chunksize = 100000
         rdf = pd.read_csv('regression_df.csv', chunksize=chunksize)

In [39]: # reducer = umap.UMAP()
         # i = 0
         # for item in rdf:
         #     if i > 3:
         #         break
         #     else:
         #         i += 1

         # data = item.drop(['Unnamed: 0', 'IDENTIFIER', 'LATEST ADMISSION DATE'], axis=1)
         # target = data[['AMER IND', 'ASIAN', 'BLACK', 'HISPANIC']]
         # data = data.drop(['AMER IND', 'ASIAN', 'BLACK', 'HISPANIC'], axis=1)
         # trans = reducer.fit_transform(data)

         # # create maskings
         # black = target.BLACK == 1
         # asian = target.ASIAN == 1
         # hisp = target.HISPANIC == 1
         # ind = target['AMER IND'] == 1
         # white = ~(black + asian + hisp)

         # plt.scatter(trans[:,0][black], trans[:,1][black], label='black', marker='.')
         # # plt.legend()
         # # plt.show()
         # plt.scatter(trans[:,0][white], trans[:,1][white], label='white', marker='.')
         # # plt.legend()
         # # plt.show()
         # plt.scatter(trans[:,0][hisp], trans[:,1][hisp], label='hisp', marker='.')
         # # plt.legend()
         # # plt.show()
         # plt.scatter(trans[:,0][ind], trans[:,1][ind], label='ind', marker='.')
         # # plt.legend()
         # # plt.show()
         # plt.scatter(trans[:,0][asian], trans[:,1][asian], label='asian', marker='.')
         # plt.legend()
         # plt.show()

In [29]: cols = ['RACE', 'GENDER', 'AGE', 'OFFENSE', 'FACILITY', 'DETAINER', 'SENTENCE DAYS']
         df = pd.read_csv('individuals.csv', usecols=cols)

In [45]: samp = df.sample(20000)
         samp.RACE = pd.factorize(samp['RACE'])[0] + 1
         samp.GENDER = pd.factorize(samp['GENDER'])[0] + 1
         samp.OFFENSE = pd.factorize(samp['OFFENSE'])[0] + 1
         samp.DETAINER = pd.factorize(samp['DETAINER'])[0] + 1
         samp.FACILITY = pd.factorize(samp['FACILITY'])[0] + 1

```

```

samp_y = samp.RACE
samp_X = samp[['GENDER', 'AGE', 'OFFENSE', 'FACILITY', 'DETAINER', 'SENTENCE DAYS']]

trans = reducer.fit_transform(samp_X)

plt.scatter(trans[:, 0], trans[:, 1], c=[sns.color_palette()[x] for x in samp_y], mar
plt.gca().set_aspect('equal', 'datalim')
plt.show()

```

/home/ethan/.local/lib/python3.7/site-packages/numba/compiler.py:602: NumbaPerformanceWarning:

The keyword argument 'parallel=True' was specified but no transformation for parallel execution

To find out why, try turning on parallel diagnostics, see <http://numba.pydata.org/numba-doc/la>

File "../.../.local/lib/python3.7/site-packages/umap/nndescent.py", line 47:

```

@numba.njit(parallel=True)
def nn_descent(
    ^


```

/home/ethan/.local/lib/python3.7/site-packages/umap/spectral.py:229: UserWarning:

Embedding a total of 50 separate connected components using meta-embedding (experimental)

/home/ethan/.local/lib/python3.7/site-packages/sklearn/manifold/spectral_embedding_.py:235: UserWarning:

Graph is not fully connected, spectral embedding may not work as expected.



It seems that UMAP doesn't work



```
In [47]: # params
ps = [5,10,20,30,40,50,60,70,80,90,100,110,120,130,140,150,160,170,180,190,200]



for p in ps:



    samp = df.sample(8000)
    samp.RACE = pd.factorize(samp['RACE'])[0] + 1
    samp.GENDER = pd.factorize(samp['GENDER'])[0] + 1
    samp.OFFENSE = pd.factorize(samp['OFFENSE'])[0] + 1
    samp.DETAINDER = pd.factorize(samp['DETAINDER'])[0] + 1
    samp.FACILITY = pd.factorize(samp['FACILITY'])[0] + 1
    samp_y = samp.RACE
    samp_X = samp[['GENDER', 'AGE', 'OFFENSE', 'FACILITY', 'DETAINDER', 'SENTENCE DAYS']]



    trans = TSNE(perplexity=p).fit_transform(samp_X)



    plt.scatter(trans[:, 0], trans[:, 1], c=[sns.color_palette()[x] for x in samp_y],
    plt.gca().set_aspect('equal', 'datalim')
    plt.title(f'p={p}')
    plt.show()
```


































```
In [70]: for off in set(samp.OFFENSE):

    mask = samp_X.OFFENSE == off

    data = samp_X[mask]

    #     trans = TSNE(perplexity=20).fit_transform(data)
    try:
        trans = reducer.fit_transform(data.values)

        plt.scatter(trans[:, 0], trans[:, 1], c=[sns.color_palette()[x] for x in samp.OFFENSE])
        plt.gca().set_aspect('equal', 'datalim')
        plt.title(f'off={off}')
        plt.show()
    except:


        continue
```

/home/ethan/.local/lib/python3.7/site-packages/umap/spectral.py:229: UserWarning:

Embedding a total of 5 separate connected components using meta-embedding (experimental)


/home/ethan/.local/lib/python3.7/site-packages/sklearn/manifold/spectral_embedding.py:235: UserWarning:

Graph is not fully connected, spectral embedding may not work as expected.




/home/ethan/.local/lib/python3.7/site-packages/umap/spectral.py:229: UserWarning:

Embedding a total of 4 separate connected components using meta-embedding (experimental)




/home/ethan/.local/lib/python3.7/site-packages/umap/spectral.py:229: UserWarning:

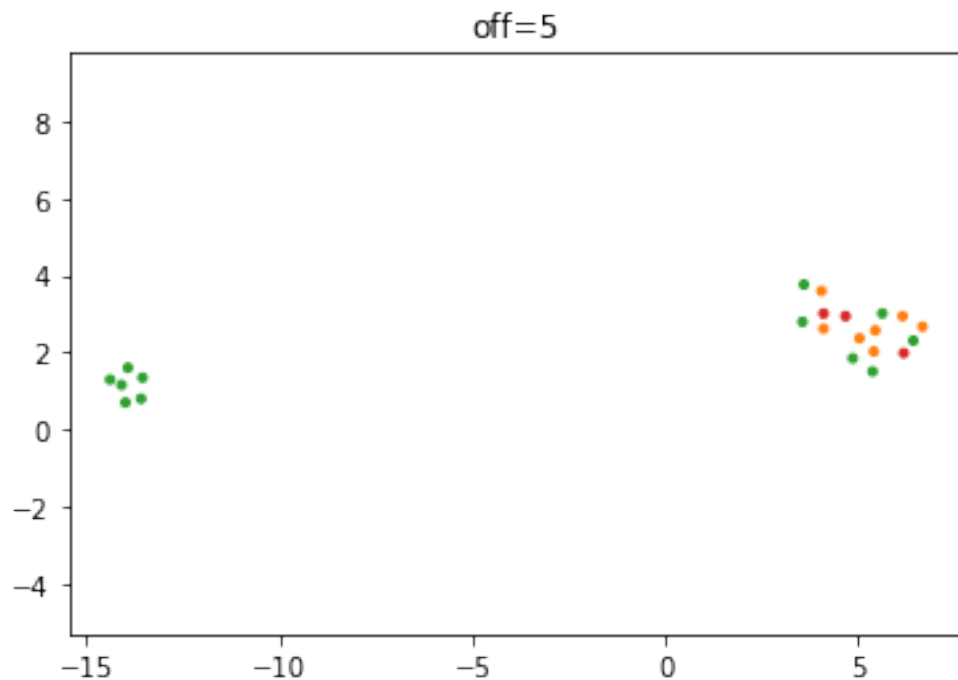
Embedding a total of 2 separate connected components using meta-embedding (experimental)



/home/ethan/.local/lib/python3.7/site-packages/umap/spectral.py:229: UserWarning:

Embedding a total of 4 separate connected components using meta-embedding (experimental)






```
/home/ethan/.local/lib/python3.7/site-packages/umap/spectral.py:229: UserWarning:
```

```
Embedding a total of 6 separate connected components using meta-embedding (experimental)
```

```
/home/ethan/.local/lib/python3.7/site-packages/sklearn/manifold/spectral_embedding_.py:235: Us
```

```
Graph is not fully connected, spectral embedding may not work as expected.
```




/home/ethan/.local/lib/python3.7/site-packages/umap/spectral.py:229: UserWarning:

Embedding a total of 8 separate connected components using meta-embedding (experimental)

/home/ethan/.local/lib/python3.7/site-packages/sklearn/manifold/spectral_embedding_.py:235: UserWarning:

Graph is not fully connected, spectral embedding may not work as expected.




```
/home/ethan/.local/lib/python3.7/site-packages/umap/spectral.py:229: UserWarning:
```

```
Embedding a total of 9 separate connected components using meta-embedding (experimental)
```

```
/home/ethan/.local/lib/python3.7/site-packages/sklearn/manifold/spectral_embedding_.py:235: UserWarning:
```

```
Graph is not fully connected, spectral embedding may not work as expected.
```



```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
```


```
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```

```
/home/ethan/.local/lib/python3.7/site-packages/scipy/sparse/linalg/eigen/arpack/arpack.py:1592
```

```
k >= N for N * N square matrix. Attempting to use scipy.linalg.eigh instead.
```


```
/home/ethan/.local/lib/python3.7/site-packages/umap/spectral.py:229: UserWarning:
```

```
Embedding a total of 3 separate connected components using meta-embedding (experimental)
```



/home/ethan/.local/lib/python3.7/site-packages/umap/spectral.py:229: UserWarning:

Embedding a total of 2 separate connected components using meta-embedding (experimental)




```
/home/ethan/.local/lib/python3.7/site-packages/umap/spectral.py:229: UserWarning:
```

```
Embedding a total of 5 separate connected components using meta-embedding (experimental)
```


```
/home/ethan/.local/lib/python3.7/site-packages/sklearn/manifold/spectral_embedding_.py:235: UserWarning:
```

```
Graph is not fully connected, spectral embedding may not work as expected.
```



```
/home/ethan/.local/lib/python3.7/site-packages/umap/spectral.py:229: UserWarning:
```

```
Embedding a total of 2 separate connected components using meta-embedding (experimental)
```





/home/ethan/.local/lib/python3.7/site-packages/umap/spectral.py:229: UserWarning:

Embedding a total of 5 separate connected components using meta-embedding (experimental)

/home/ethan/.local/lib/python3.7/site-packages/sklearn/manifold/spectral_embedding_.py:235: UserWarning:

Graph is not fully connected, spectral embedding may not work as expected.




```
/home/ethan/.local/lib/python3.7/site-packages/umap/spectral.py:229: UserWarning:
```

```
Embedding a total of 7 separate connected components using meta-embedding (experimental)
```


```
/home/ethan/.local/lib/python3.7/site-packages/sklearn/manifold/spectral_embedding_.py:235: UserWarning:
```


```
Graph is not fully connected, spectral embedding may not work as expected.
```



/home/ethan/.local/lib/python3.7/site-packages/umap/spectral.py:229: UserWarning:

Embedding a total of 3 separate connected components using meta-embedding (experimental)







/home/ethan/.local/lib/python3.7/site-packages/umap/spectral.py:229: UserWarning:

Embedding a total of 8 separate connected components using meta-embedding (experimental)



/home/ethan/.local/lib/python3.7/site-packages/sklearn/manifold/spectral_embedding_.py:235: UserWarning:



Graph is not fully connected, spectral embedding may not work as expected.



/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:


`n_neighbors` is larger than the dataset size; truncating to `X.shape[0] - 1`






/home/ethan/.local/lib/python3.7/site-packages/umap/spectral.py:229: UserWarning:


Embedding a total of 4 separate connected components using meta-embedding (experimental)



```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```







/home/ethan/.local/lib/python3.7/site-packages/umap/spectral.py:229: UserWarning:
Embedding a total of 4 separate connected components using meta-embedding (experimental)







/home/ethan/.local/lib/python3.7/site-packages/umap/spectral.py:229: UserWarning:


Embedding a total of 4 separate connected components using meta-embedding (experimental)
















/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1









/home/ethan/.local/lib/python3.7/site-packages/umap/spectral.py:229: UserWarning:

Embedding a total of 2 separate connected components using meta-embedding (experimental)





```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```





```
/home/ethan/.local/lib/python3.7/site-packages/umap/spectral.py:229: UserWarning:
```

```
Embedding a total of 8 separate connected components using meta-embedding (experimental)
```

```
/home/ethan/.local/lib/python3.7/site-packages/sklearn/manifold/spectral_embedding_.py:235: UserWarning:
```


```
Graph is not fully connected, spectral embedding may not work as expected.
```






/home/ethan/.local/lib/python3.7/site-packages/umap/spectral.py:229: UserWarning:


Embedding a total of 3 separate connected components using meta-embedding (experimental)



```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```



```
/home/ethan/.local/lib/python3.7/site-packages/umap/spectral.py:229: UserWarning:  
Embedding a total of 3 separate connected components using meta-embedding (experimental)
```






```
/home/ethan/.local/lib/python3.7/site-packages/umap/spectral.py:229: UserWarning:
```

Embedding a total of 6 separate connected components using meta-embedding (experimental)


```
/home/ethan/.local/lib/python3.7/site-packages/sklearn/manifold/spectral_embedding.py:235: Us
```


Graph is not fully connected, spectral embedding may not work as expected.






/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1






/home/ethan/.local/lib/python3.7/site-packages/umap/spectral.py:229: UserWarning:


Embedding a total of 3 separate connected components using meta-embedding (experimental)





/home/ethan/.local/lib/python3.7/site-packages/umap/spectral.py:229: UserWarning:

Embedding a total of 2 separate connected components using meta-embedding (experimental)





```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```



/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:

n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1

/home/ethan/.local/lib/python3.7/site-packages/scipy/sparse/linalg/eigen/arpack/arpack.py:1592

k >= N for N * N square matrix. Attempting to use scipy.linalg.eigh instead.

/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:


n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1



/home/ethan/.local/lib/python3.7/site-packages/scipy/sparse/linalg/eigen/arpack/arpack.py:1592


k >= N for N * N square matrix. Attempting to use scipy.linalg.eigh instead.

/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:


n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1










/home/ethan/.local/lib/python3.7/site-packages/umap/spectral.py:229: UserWarning:
Embedding a total of 3 separate connected components using meta-embedding (experimental)



```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```




```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```



/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:



`n_neighbors` is larger than the dataset size; truncating to `X.shape[0] - 1`



`/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:`



`n_neighbors` is larger than the dataset size; truncating to `X.shape[0] - 1`







/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:


`n_neighbors` is larger than the dataset size; truncating to `X.shape[0] - 1`




```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```




```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```




```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```




/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1






```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```






/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1







```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```




/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1




```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```





```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```



```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```






```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```


```
/home/ethan/.local/lib/python3.7/site-packages/scipy/sparse/linalg/eigen/arpack/arpack.py:1592  
k >= N for N * N square matrix. Attempting to use scipy.linalg.eigh instead.
```

```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```




```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```


```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```


/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1



```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```



```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```




```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
```

n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1


```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
```

n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1




```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```

```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```



```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```



```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
```


```
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```

```
/home/ethan/.local/lib/python3.7/site-packages/scipy/sparse/linalg/eigen/arpack/arpack.py:1592
```

```
k >= N for N * N square matrix. Attempting to use scipy.linalg.eigh instead.
```


```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
```

```
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```



```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
```

```
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```



```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
```

```
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```

```
/home/ethan/.local/lib/python3.7/site-packages/scipy/sparse/linalg/eigen/arpack/arpack.py:1592
```

```
k >= N for N * N square matrix. Attempting to use scipy.linalg.eigh instead.
```

```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
```


```
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```

```
/home/ethan/.local/lib/python3.7/site-packages/scipy/sparse/linalg/eigen/arpack/arpack.py:1592
```


```
k >= N for N * N square matrix. Attempting to use scipy.linalg.eigh instead.
```

```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
```


```
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```




/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1





```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```



```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```



```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```




```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
```

```
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```


```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
```

```
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```



```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
```

```
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```




```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
```


```
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```

```
/home/ethan/.local/lib/python3.7/site-packages/scipy/sparse/linalg/eigen/arpack/arpack.py:1592
```


```
k >= N for N * N square matrix. Attempting to use scipy.linalg.eigh instead.
```




/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1



```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```





```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
```

```
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```

```
/home/ethan/.local/lib/python3.7/site-packages/scipy/sparse/linalg/eigen/arpack/arpack.py:1592
```


```
k >= N for N * N square matrix. Attempting to use scipy.linalg.eigh instead.
```

```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
```

```
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```

```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
```

```
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```



```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
```

```
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```

```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
```

```
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```

```
/home/ethan/.local/lib/python3.7/site-packages/scipy/sparse/linalg/eigen/arpack/arpack.py:1592
```

```
k >= N for N * N square matrix. Attempting to use scipy.linalg.eigh instead.
```

```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
```


```
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```

```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
```

```
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```

```
/home/ethan/.local/lib/python3.7/site-packages/scipy/sparse/linalg/eigen/arpack/arpack.py:1592
```

```
k >= N for N * N square matrix. Attempting to use scipy.linalg.eigh instead.
```

```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
```

```
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```

```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
```

```
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```

```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
```

```
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```

```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
```



```
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```

```
/home/ethan/.local/lib/python3.7/site-packages/scipy/sparse/linalg/eigen/arpack/arpack.py:1592
```

```
k >= N for N * N square matrix. Attempting to use scipy.linalg.eigh instead.
```


```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
```

```
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```





/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:

`n_neighbors` is larger than the dataset size; truncating to `X.shape[0] - 1`




`/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:`

`n_neighbors` is larger than the dataset size; truncating to `X.shape[0] - 1`



/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:

n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1



```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
```

```
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```

```
/home/ethan/.local/lib/python3.7/site-packages/scipy/sparse/linalg/eigen/arpack/arpack.py:1592
```

```
k >= N for N * N square matrix. Attempting to use scipy.linalg.eigh instead.
```

```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
```



```
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```



```
/home/ethan/.local/lib/python3.7/site-packages/scipy/sparse/linalg/eigen/arpack/arpack.py:1592
```



```
k >= N for N * N square matrix. Attempting to use scipy.linalg.eigh instead.
```

```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:
```

```
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```







/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:


n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1


/home/ethan/.local/lib/python3.7/site-packages/scipy/sparse/linalg/eigen/arpack/arpack.py:1592

k >= N for N * N square matrix. Attempting to use scipy.linalg.eigh instead.


/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:

n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1








```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```







```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```







```
/home/ethan/.local/lib/python3.7/site-packages/umap/umap_.py:1383: UserWarning:  
n_neighbors is larger than the dataset size; truncating to X.shape[0] - 1
```




































```
In [73]: for off in set(samp.OFFENSE):  
  
    mask = samp_X.OFFENSE == off  
  
    data = samp_X[mask]  
  
    try:  
        trans = TSNE(perplexity=12).fit_transform(data)  
  
        plt.scatter(trans[:, 0], trans[:, 1], c=[sns.color_palette()[x] for x in samp.OFFENSE])  
        plt.gca().set_aspect('equal', 'datalim')  
        plt.title(f'off={off}')  
        plt.show()  
    except:  
  
        continue
```















































































































































































































































































TSNE is having mized results while UMAP doesn't seem to capture any good groupings. Maybe I should go straight to kmeans or kdtrees\ I should block by crime and examine a few "big" crimes the try several different perplexities.

```
In [ ]:
```